

Introduction to Machine Learning

Lecture 6: Gradient boosting

Harbour.Space University
February 2021

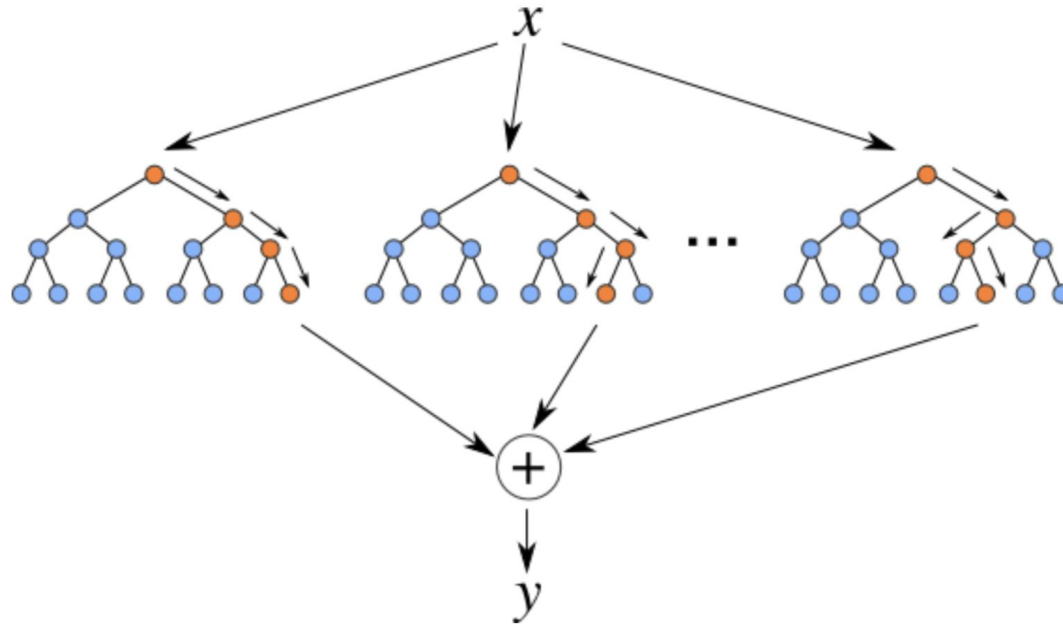
Nikolay Karpachev

Outline

1. Random Forest recap. Out-of-bag error.
2. Boosting technique
3. Gradient Boosting

Random Forest

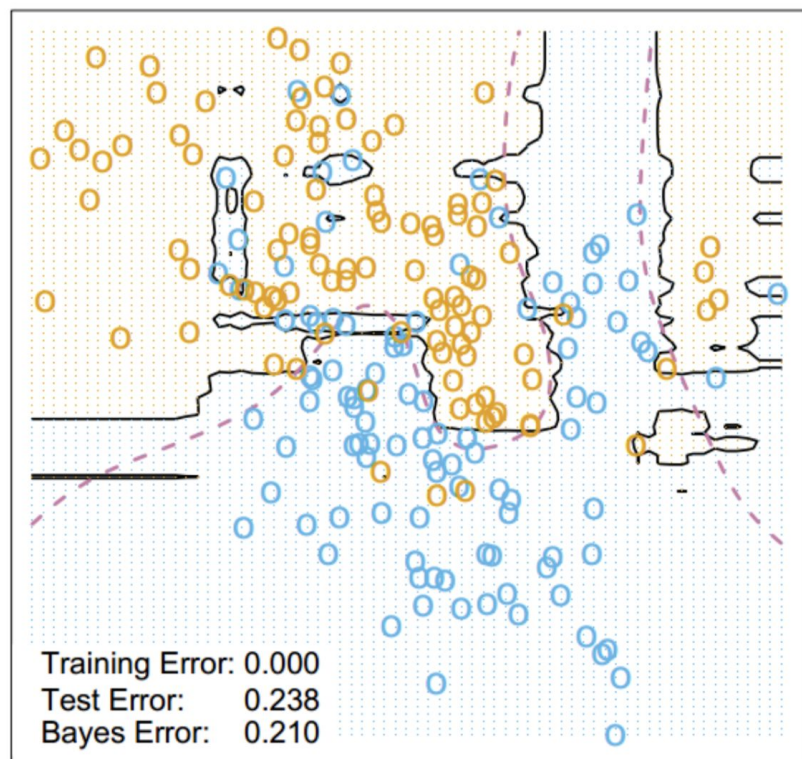
Bagging + RSM = Random Forest



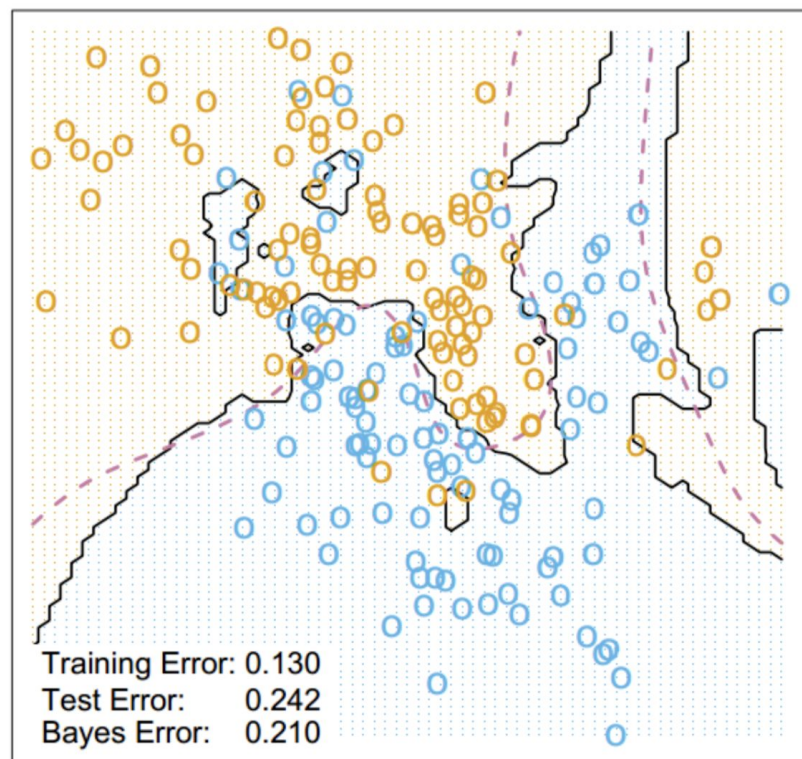
- One of the greatest “universal” models.
- There are some modifications: Extremely Randomized Trees, Isolation Forest, etc.
- Allows to use train data for validation: OOB

$$\text{OOB} = \sum_{i=1}^{\ell} L \left(y_i, \frac{1}{\sum_{n=1}^N [x_i \notin X_n]} \sum_{n=1}^N [x_i \notin X_n] b_n(x_i) \right)$$


Random Forest Classifier




3-Nearest Neighbors



Bias-variance tradeoff

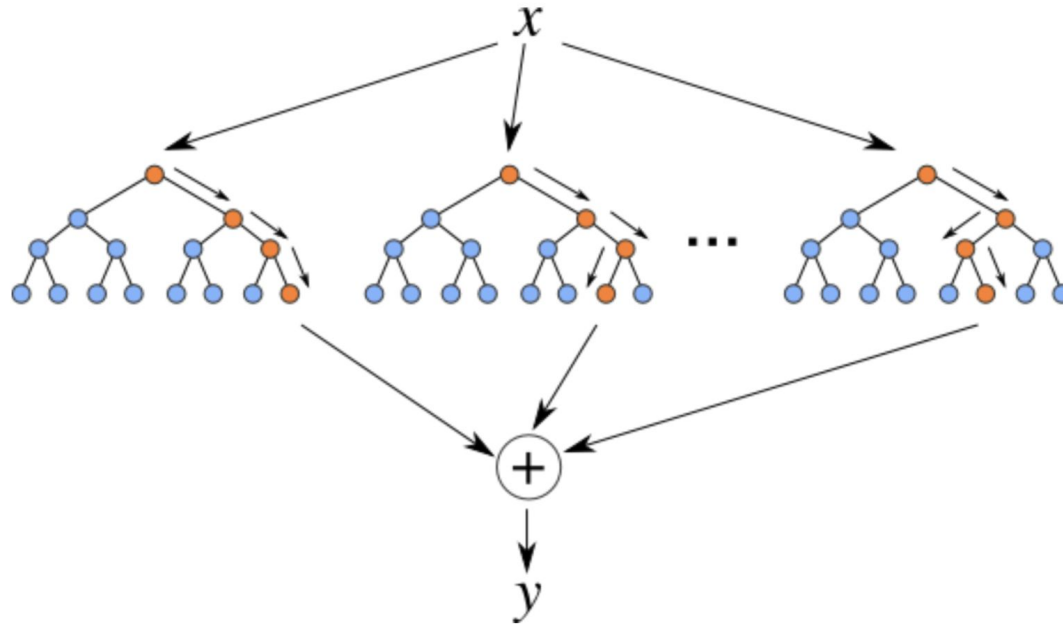


Bias-variance tradeoff



Random Forest

Is Random Forest decreasing bias or variance by building the trees ensemble?



Boosting

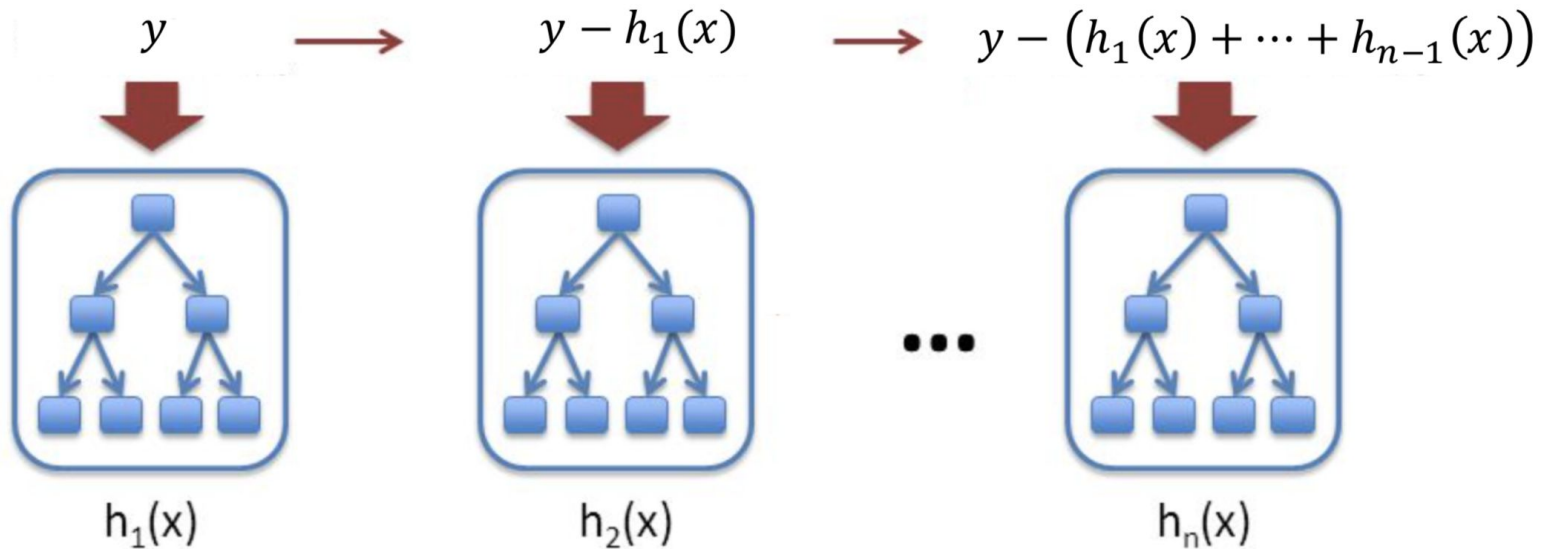
Boosting is a technique to combine multiple weak algorithms into one strong ensemble.

- Boosting adds algorithms incrementally
- At the current step, ensemble consists of N algorithms
- Algorithm $N+1$ is trained with respect to already built composition to minimize the total error

Boosting


Example: regression

$$a_n(x) = h_1(x) + \dots + h_n(x)$$


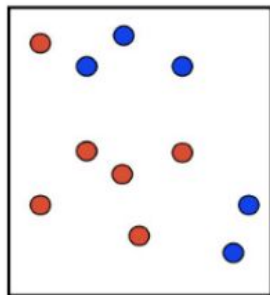


Boosting


Binary classification problem.
Models - decision stumps.



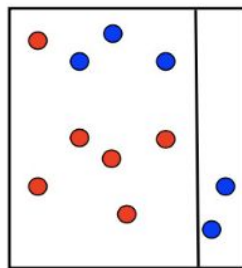
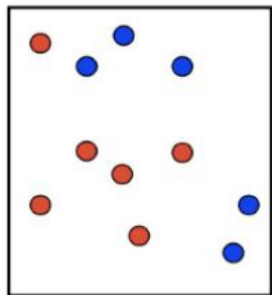
Boosting: intuition




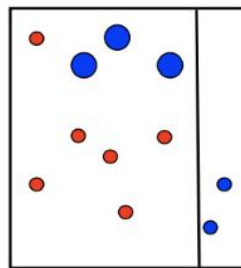
$t = 1$




Boosting: intuition





$t = 1$





$t = 2$





Boosting: intuition




$t = 1$



$t = 2$





$t = 3$



Boosting: intuition

Binary classification problem.
Models - decision stumps.



Gradient boosting: theory

Denote dataset $\{(x_i, y_i)\}_{i=1, \dots, n}$, loss function $L(y, f)$.

Gradient boosting: theory

Denote dataset $\{(x_i, y_i)\}_{i=1, \dots, n}$, loss function $L(y, f)$.

Optimal model:

$$\hat{f}(x) = \arg \min_{f(x)} L(y, f(x)) = \arg \min_{f(x)} \mathbb{E}_{x,y}[L(y, f(x))]$$

Gradient boosting: theory

Denote dataset $\{(x_i, y_i)\}_{i=1, \dots, n}$, loss function $L(y, f)$.

Optimal model:

$$\hat{f}(x) = \arg \min_{f(x)} L(y, f(x)) = \arg \min_{f(x)} \mathbb{E}_{x,y}[L(y, f(x))]$$

Let it be from parametric family: $\hat{f}(x) = f(x, \hat{\theta})$,

$$\hat{\theta} = \arg \min_{\theta} \mathbb{E}_{x,y}[L(y, f(x, \theta))]$$

Gradient boosting: theory

- Let's assume real-valued y
- Gradient boosting builds a single model as a weighted sum of “*base learners*”

$$\hat{F}(x) = \sum_{i=1}^M \gamma_i h_i(x) + \text{const.}$$

Gradient boosting: theory

- The goal is to minimize the loss function on the training set

Boosting procedure:

1. Start with the constant function

$$F_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma),$$

2. Incrementally expand the composition by minimizing the **residual error**

$$F_m(x) = F_{m-1}(x) + \arg \min_{h_m \in \mathcal{H}} \left[\sum_{i=1}^n L(y_i, F_{m-1}(x_i) + h_m(x_i)) \right],$$

↙
We select the
optimal h_m

↘
Loss of the new composition

Time for your questions and a coffee
break

Gradient boosting: theory

$$F_m(x) = F_{m-1}(x) + \arg \min_{h_m \in \mathcal{H}} \left[\sum_{i=1}^n L(y_i, F_{m-1}(x_i) + h_m(x_i)) \right],$$

Diagram illustrating the equation for $F_m(x)$ in gradient boosting:

- An arrow points from $\arg \min_{h_m \in \mathcal{H}}$ to the text: "We select the optimal h_m ".
- An arrow points from $L(y_i, F_{m-1}(x_i) + h_m(x_i))$ to the text: "Loss of the new composition".


How to find the optimal h_m ?

Gradient boosting: theory

Key idea: apply gradient descent in the space of model predictions

$$F_m(x) = F_{m-1}(x) - \gamma_m \sum_{i=1}^n \nabla_{F_{m-1}} L(y_i, F_{m-1}(x_i)),$$



Loss gradient at current prediction point



$$\gamma_m = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, F_{m-1}(x_i) - \gamma \nabla_{F_{m-1}} L(y_i, F_{m-1}(x_i))),$$

“Learning rate”

Find the optimal “learning rate”



- We fit the next base learner to the anti-gradient of the error
- Thus, the composition resembles gradient descent
- Each base learner addition = gradient step in functional space

Gradient boosting: theory

Complete algorithm:

- Initialize the model with a constant value

$$F_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma).$$

Incrementally:

- Compute *pseudo-residuals* for each data point

$$r_{im} = - \left[\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} \right]_{F(x)=F_{m-1}(x)} \quad \text{for } i = 1, \dots, n.$$

- Fit new base learner to pseudo-residuals
- Find the optimal “learning rate” coefficient
- Update the model

$$\gamma_m = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, F_{m-1}(x_i) + \gamma h_m(x_i)).$$

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x).$$

Gradient boosting: beautiful demo

Great demo:

http://arogozhnikov.github.io/2016/06/24/gradient_boosting_explained.html

Gradient boosting: theory

What we need:

- Data.
- Loss function and its gradient.
- Family of algorithms (with constraints on hyperparameters if necessary).
- Number of iterations M .
- Initial value (GBM by Friedman): constant.

Gradient boosting: example


What we need:

- Data: toy dataset $y = \cos(x) + \epsilon, \epsilon \sim \mathcal{N}(0, \frac{1}{5}), x \in [-5, 5]$
- Loss function: MSE
- Family of algorithms: decision trees with depth 2
- Number of iterations $M = 3$
- Initial value: just mean value


Gradient boosting: example

Left: full ensemble on each step.


Right: additional tree decisions.



Spam Data



California Housing Data



Gradient boosting: regularization

Boosting is a very powerful algorithm

But it can overfit the training data and lose the generalization ability

We need to regularize the model

Regularization options:

- Constrain number of trees
- Set max. depth of the trees

Thanks for the attention!